

MATH562 - Theory of Machine Learning

Based on lectures from Winter 2026 by Prof. Courtney Paquette.
Notes by Louis Meunier

Contents

1	Introduction	2
1.1	Some Linear Algebra	2
1.1.1	Inverting Block Matrices	2
1.1.2	Eigenvalues and Singular Values	2
1.2	Concentration Inequalities	2
1.2.1	Hoeffding Inequality	3
1.2.2	McDiarmid's Inequality	6
1.2.3	Bernstein's Inequality	6
1.2.4	Expectation of the Maximum	7
2	Introduction to Supervised Learning	7
2.1	Training Data Predictions	7
2.2	Decision Theory	8
2.2.1	Supervised Learning and Loss Functions	8
2.2.2	Risks	8
2.2.3	Baye's Risk, Predictor	9
2.3	Empirical Risk Minimization	10
2.3.1	Risk Decomposition	11
2.4	Statistical Learning Theory	12
2.4.1	Measures of Performance	12
2.4.2	Notions of Consistency over Classes of Probability distributions	12

§1 INTRODUCTION

§1.1 Some Linear Algebra

1.1.1 Inverting Block Matrices

Let

$$M := \begin{pmatrix} A & B \\ C & D \end{pmatrix} \in \mathbb{R}^{(p+q) \times (p+q)},$$

i.e. $A \in \mathbb{R}^{p \times p}$, $B \in \mathbb{R}^{p \times q}$, $C \in \mathbb{R}^{q \times p}$ and $D \in \mathbb{R}^{q \times q}$ (where we use the convention that if $A \in \mathbb{R}^{m \times n}$, then A has m rows and n columns, so in particular maps $\mathbb{R}^n \rightarrow \mathbb{R}^m$). If A is invertible, let

$$M \setminus A := D - CA^{-1}B =: \text{Schur Complement (of } A \text{ with respect to } M\text{)}.$$

Then,

$$M^{-1} = \begin{pmatrix} A^{-1} + A^{-1}B(M \setminus A)^{-1}CA^{-1} & -A^{-1}B(M \setminus A)^{-1} \\ -(M \setminus A)^{-1}CA^{-1} & (M \setminus A)^{-1} \end{pmatrix}.$$

Similarly, if D invertible and $M \setminus D := A - BD^{-1}C$, then

$$M^{-1} = \begin{pmatrix} (M \setminus D)^{-1} & -(M \setminus D)^{-1}BD^{-1} \\ -D^{-1}C(M \setminus D)^{-1} & D^{-1} + D^{-1}C(M \setminus D)^{-1}BD^{-1} \end{pmatrix}.$$

1.1.2 Eigenvalues and Singular Values

Given $A \in \mathbb{R}^{n \times n}$ symmetric, there exists $U \in \mathbb{R}^{n \times n}$ orthogonal (i.e., $U^T = U^{-1}$) such that

$$A = U \text{ diag}(\lambda) U^T,$$

where $\lambda = (\lambda_1, \dots, \lambda_n)$ for λ_i 's the eigenvectors of A . In particular, if $U^{(i)}$ enumerate the columns of U , we have

$$AU^{(i)} = \lambda_i U^{(i)},$$

i.e. the $U^{(i)}$'s are the eigenvectors of A .

Given $X \in \mathbb{R}^{n \times d}$, $n \geq d$, then there exists an orthogonal matrix $V \in \mathbb{R}^{d \times d}$ and $U \in \mathbb{R}^{n \times d}$ with orthogonal columns, and a matrix of *singular values* $s \in \mathbb{R}_+^d = \{(v_1, \dots, v_d) \in \mathbb{R}^d \mid v_i \geq 0 \forall i = 1, \dots, d\}$ such that

$$X = U \text{ Diag}(s) V^T.$$

Remark 1.1:

1. if $u_i \in \mathbb{R}^n$, $v_j \in \mathbb{R}^d$ are the columns of U , V resp., then $X = \sum_{i=1}^d s_i u_i v_i^T$
2. if s_i a singular value of X , then s_i^2 an eigenvalue of XX^T and $X^T X$.

§1.2 Concentration Inequalities

Here we study the question of how the magnitude of the average of n independent, mean 0 random variables behaves compared to their average magnitude, specifically with respect to n .

We know that the central-limit theorem states that for z_i iid with variance σ^2 , $\sqrt{n}(\frac{1}{n} \sum z_i - \mathbb{E}[z])$ converges in distribution to a $\mathcal{N}(0, \sigma^2)$; this is an asymptotic result, which gives no

information about the rate of this converge with respect to n , which is what we care about in our study.

→**Proposition 1.1** (Markov's): Let Y be a nonnegative r.v. with finite mean. Then,

$$\mathbb{P}(Y \geq \varepsilon) \leq \frac{1}{\varepsilon} \mathbb{E}[Y], \quad \forall \varepsilon > 0.$$

→**Proposition 1.2** (Chebyshev's): Let X be a r.v. with finite mean and variance, then

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq \varepsilon) \leq \frac{\text{Var}[X]}{\varepsilon^2}, \quad \forall \varepsilon > 0.$$

→**Corollary 1.1**: If $z_i, i = 1, \dots, n$ are iid with variance σ^2 , then

$$\mathbb{P}\left(\left|\frac{1}{n} \sum_{i=1}^n z_i - \mathbb{E}[z]\right| \geq \varepsilon\right) \leq \frac{\sigma^2}{n\varepsilon^2}.$$

→**Proposition 1.3** (Union Bound, Max/Tail Bound):

1. $\mathbb{P}\left(\bigcup_{f \in \mathcal{F}} A_f\right) \leq \sum_{f \in \mathcal{F}} \mathbb{P}(A_f)$
2. $\mathbb{P}\left(\sup_{f \in \mathcal{F}} Z_f \geq t\right) \leq \sum_{f \in \mathcal{F}} \mathbb{P}(Z_f \geq t)$

→**Proposition 1.4** (Jensen's Inequality): If $F : \mathbb{R} \rightarrow \mathbb{R}$ convex and X an r.v. with finite mean,

$$F(\mathbb{E}[X]) \leq \mathbb{E}[F(X)].$$

1.2.1 Hoeffding Inequality

→**Proposition 1.5** (Hoeffding Inequality): Let z_1, \dots, z_n be independent r.v.s with $z_i \in [0, 1]$ a.s.. Then, for any $t \geq 0$,

$$\mathbb{P}\left(\frac{1}{n} \sum_{i=1}^n z_i - \frac{1}{n} \sum_{i=1}^n \mathbb{E}[z_i] \geq t\right) \leq \exp(-2nt^2).$$

Remark 1.2: Read this result as a *fast* (exponential) convergence of the empirical mean to the true mean as the sample size n grows.

PROOF. First we claim that

$$z \in [0, 1] \text{ a.s.} \Rightarrow \mathbb{E}[\exp(s(z - \mathbb{E}[z]))] \leq \exp\left(\frac{s^2}{8}\right). \quad (\dagger)$$

We'll assume z is centered for the sake of notation. Let $\varphi(s) := \log(\mathbb{E}[\exp(sz)])$.

Remark that

$$\begin{aligned}\varphi(0) &= 0 \\ \varphi'(s) &= \frac{\mathbb{E}[z \exp(sz)]}{\mathbb{E}[\exp(sz)]} \\ \varphi''(s) &= \frac{\mathbb{E}[z^2 \exp(sz)]}{\mathbb{E}[\exp(sz)]} - \left(\frac{\mathbb{E}[z \exp(sz)]}{\mathbb{E}[\exp(sz)]} \right)^2.\end{aligned}$$

In particular, if we define a new probability density

$$\tilde{z} \mapsto \frac{e^{s\tilde{z}}}{\mathbb{E}[e^{sz}]}$$

with respect to that of z , and let \tilde{z} be distributed with respect to this distribution, then

$$\text{Var}(\tilde{z}) = \varphi''(s).$$

Note that $\tilde{z} \in [0, 1]$ a.s.. In addition, we have that

$$\begin{aligned}\text{Var}(\tilde{z}) &= \inf_{v \in [0, 1]} \mathbb{E}[(\tilde{z} - v)^2] \\ &\leq \mathbb{E}\left[\left(\tilde{z} - \frac{1}{2}\right)^2\right] = \frac{1}{4} \mathbb{E}\left[\left(\underbrace{\frac{2\tilde{z}-1}{\leq 1 \text{ a.s.}}}_{\leq 1 \text{ a.s.}}\right)^2\right] \leq \frac{1}{4},\end{aligned}$$

so that $\varphi''(s) \leq \frac{1}{4}$ for all s . Thus, by Taylor expanding φ , we find

$$\varphi(s) \leq \varphi(0) + \varphi'(0)s + \frac{s^2}{2} \frac{1}{4} = \frac{s^2}{8},$$

using the bound above and the fact $\varphi'(0) = 0$ (checking the above formula). Thus,

$$\varphi(s) = \log(\mathbb{E}[\exp(sz)]) \leq \frac{s^2}{8},$$

from which the claim (\dagger) follows by taking \exp of both sides.

Next, let $t \geq 0$ and put $\bar{z} = \frac{1}{n} \sum z_i$. Then,

$$\begin{aligned}\mathbb{P}(\bar{z} - \mathbb{E}[\bar{z}] \geq t) &= \mathbb{P}(\exp(s(\bar{z} - \mathbb{E}[\bar{z}])) \geq \exp(st)) \\ (\text{Markov's}) \quad &\leq e^{-st} \mathbb{E}[\exp(s(\bar{z} - \mathbb{E}[\bar{z}]))] \\ (\text{Indep.}) \quad &= e^{-st} \prod_{i=1}^n \mathbb{E}\left[\exp\left(\frac{s}{n}(z_i - \mathbb{E}[z_i])\right)\right] \\ (\dagger) \quad &\leq e^{-st} \prod_{i=1}^n \exp\left(\frac{s^2}{8n^2}\right) = \exp\left(-st + \frac{s^2}{8n}\right).\end{aligned}$$

This bound held for all s , so in particular holds at $\bar{s} = \arg\min\left\{-st + \frac{s^2}{8n}\right\} = 4nt$. Plugging in this value for s gives the result. ■

→**Corollary 1.2** (2-sided Hoeffding): With the same hypotheses as the previous proposition, we have

$$\mathbb{P}\left(\left|\frac{1}{n} \sum z_i - \frac{1}{n} \sum \mathbb{E}[z_i]\right| \geq t\right) \leq 2 \exp(-2nt^2), \forall t \geq 0.$$

If instead $z_i \in [a, b]$ a.s., we can replace the rhs with

$$\leq 2 \exp\left(\frac{-2nt^2}{(a-b)^2}\right).$$

Remark 1.3:

- How is Hoeffding used? Start with a probability, then derive the necessary t (usually, as a function of n) to achieve that bound. e.g., if $z_i \in [a, b]$ a.s. and for any $\delta \in (0, 1)$, then with probability $1 - \delta$,

$$\left|\frac{1}{n} \sum z_i - \frac{1}{n} \sum \mathbb{E}[z_i]\right| < \frac{|a-b|}{\sqrt{2n}} \sqrt{\log\left(\frac{2}{\delta}\right)}$$

- An extension exists for martingales. If $Z_i, i = 1, \dots, n$ martingales with respect to a filtration $\{\mathcal{F}_i\}$ and $|Z_i| \leq c_i$ a.s., then

$$\mathbb{P}\left(\frac{1}{n} \sum_{i=1}^n Z_i \geq t\right) \leq \exp\left(-\frac{n^2 t^2}{2\|c\|^2}\right), \quad c := (c_1, \dots, c_n).$$

→**Definition 1.1** (Sub-Gaussian): We say an r.v. X is *sub-Gaussian* if there exists $\tau \in \mathbb{R}_+$ such that

$$\mathbb{E}[\exp(s(X - \mathbb{E}[X]))] \leq \exp(\tau^2 s^2), \quad \forall s \in \mathbb{R}.$$

We define the *sub-Gaussian norm* by

$$\|X\|_{\psi_2} := \inf\left\{k \geq 0 : \mathbb{E}\left[\exp\left(\frac{X^2}{k^2}\right)\right] \leq 2\right\},$$

i.e. the “best” sub-Gaussian parameter for X .

Remark 1.4: Interpretation: X has tails decaying as fast (or faster) than a Gaussian.

→**Proposition 1.6:** X is sub-Gaussian iff there exists a $k \in \mathbb{R}_+$ such that

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq t) \leq 2 \exp\left(-\frac{t^2}{k^2}\right), \quad \forall t \in \mathbb{R}.$$

→ **Definition 1.2** (Sub-Exponential): We say X sub-exponential if

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq t) \leq 2 \exp\left(-\frac{t}{k}\right),$$

for some k and for all $t \geq 0$. We define the *sub-Gaussian norm* by

$$\|X\|_{\psi_1} := \inf\left\{k \geq 0 : \mathbb{E}\left[\exp\left(\frac{|X|}{k^2}\right)\right] \leq 2\right\},$$

i.e. the “best” sub-Gaussian parameter for X .

Remark 1.5: This is a similar, but slower, tail bound than sub-Gaussian.

1.2.2 McDiarmid's Inequality

For a measure space Z and nonnegative integer n , we say $f : Z^n \rightarrow \mathbb{R}$ is a function of bounded variation with constant c if for all $i \in [n] := \{1, \dots, n\}$ and points $z_1, \dots, z_n, z'_i \in Z$, then

$$|f(z_1, \dots, z_i, \dots, z_n) - f(z_1, \dots, z'_i, \dots, z_n)| \leq c.$$

→ **Proposition 1.7** (McDiarmid's Inequality): Let z_1, \dots, z_n be independent r.v.s on some measure space Z and $f : Z^n \rightarrow \mathbb{R}$ be a function of bounded variation with constant c . Then,

$$\mathbb{P}(|f(z_1, \dots, z_n) - \mathbb{E}[f(z_1, \dots, z_n)]| \geq t) \leq 2 \exp\left(-\frac{2t^2}{nc^2}\right), \quad \forall t \geq 0.$$

Remark 1.6: We can extend this to $z_i \in \mathbb{R}^d$; if $\|z_i\|_2 \leq c$ a.s., then $\left\|\frac{1}{n} \sum z_i\right\|_2 \leq \frac{c}{\sqrt{n}} \left(1 + \sqrt{2 \log\left(\frac{1}{\delta}\right)}\right)$ with probability $\geq 1 - \delta$.

1.2.3 Bernstein's Inequality

→ **Proposition 1.8** (Bernstein's): Let $z_i, i = 1, \dots, n$ be independent with $|z_i| \leq c$ a.s. and mean zero. Then for all $t \geq 0$,

$$\mathbb{P}\left(\frac{1}{n} \left| \sum_i z_i \right| \geq t\right) \leq 2 \exp\left(-\frac{nt^2}{2\sigma^2 + 2ct/3}\right), \quad \sigma^2 := \frac{1}{n} \sum_i \text{Var}(z_i).$$

In particular, for $\delta \in (0, 1)$, then with probability $\geq 1 - \delta$,

$$\left| \frac{1}{n} \sum z_i \right| \leq \sqrt{\frac{2\sigma^2 \log(\frac{2}{\delta})}{n}} + \frac{2c \log(\frac{2}{\delta})}{3n}.$$

→ **Proposition 1.9** (Extension of Bernstein's, sub-exponential): Let x_1, \dots, x_n be mean zero, independent, sub-exponential r.v.s with constants k_i , and let $a \in \mathbb{R}^n$. Then, for all $t \geq 0$,

$$\mathbb{P}(|\sum a_i x_i| \geq t) \leq 2 \exp\left(-c \min\left\{\frac{t^2}{k^2 \|a\|_2^2}, \frac{1}{k \|a\|_\infty}\right\}\right).$$

→ **Proposition 1.10** (Extension of Bernstein's, non-zero means): With the same hypothesis as Bernstein's but without the zero means, we have

$$\mathbb{P}\left(\left|\frac{1}{n} \sum z_i - \frac{1}{n} \sum \mathbb{E}[z_i]\right| \geq t\right) \leq 2 \exp\left(-\frac{nt^2}{2\sigma^2 + 2c\frac{t}{3}}\right).$$

1.2.4 Expectation of the Maximum

→ **Proposition 1.11:** Let z_i be (possibly dependent) mean-zero, \mathbb{R} -values r.v.s which are all sub-Gaussian with constant τ^2 . Then,

$$\mathbb{E}[\max\{z_1, \dots, z_n\}] \leq \sqrt{2\tau^2 \log(n)}.$$

PROOF. For all $t > 0$,

$$\begin{aligned} \mathbb{E}[\max\{z_1, \dots, z_n\}] &\leq \frac{1}{t} \log(\mathbb{E}[\exp(t \max(z_i))]) \quad (\text{Jensen's}) \\ &= \frac{1}{t} \log(\mathbb{E}[\max\{\exp(tZ_i)\}]) \quad (\exp \text{ increasing}) \\ &\leq \frac{1}{t} \log(\mathbb{E}[\sum \exp(tZ_i)]) \quad (\max \text{ leq sum}) \\ &\leq \frac{1}{t} \log\left(n \exp\left(\tau^2 \frac{t^2}{2}\right)\right) \quad (\text{sub-Gaussian}) \\ &= \frac{\log(n)}{t} + \frac{\tau^2 t}{2}. \end{aligned}$$

The proof follows by taking $t := \tau^{-1} \sqrt{2 \log(n)}$. ■

§2 INTRODUCTION TO SUPERVISED LEARNING

§2.1 Training Data Predictions

The goal of supervised learning is to take a series of observations $(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}$ for $i \in [n]$ (called *training data*) and to predict a new $y \in \mathcal{Y}$ given a (previously unseen) $x \in \mathcal{X}$ (*testing data*).

We write

- \mathcal{X} for our space of *inputs*, typically embedded in \mathbb{R}^d (where d tends to be very large; think images encoded as large matrices of pixels, text, videos, etc)
- \mathcal{Y} for our space of *outputs* or *labels* for the data

The challenges in supervised learning are twofold:

1. $y \in \mathcal{Y}$ may not be a deterministic function of $x \in \mathcal{X}$
2. inputs may live in a high-dimensional space, hence it is computationally expensive to work with them

We make two primary blanket assumptions of our problem:

1. we aim to maximize the expectation of some measure of performance with respect to some testing distribution we put on our data

2. we assume (x_i, y_i) are iid, with the training data having the same distribution as the testing data

→**Definition 2.1** (Machine Learning (ML) Algorithm): An *ML algorithm* is a function from the data set, $(\mathcal{X} \times \mathcal{Y})^n$ to a function $\mathcal{X} \rightarrow \mathcal{Y}$.

§2.2 Decision Theory

The question we aim to answer here is, what is the *optimal* performance of an algorithm, regardless of the finiteness of the data? I.e., if we havd perfect knowledge of the underlying probability distribution of our data, how should we design our algorithm?

We assume for now a fixed (testing) distribution $P_{x,y}$ on $\mathcal{X} \times \mathcal{Y}$ with P_x marginal distribution on \mathcal{X} .

2.2.1 Supervised Learning and Loss Functions

→**Definition 2.2** (Loss Function): A *loss function* is a mapping $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ where $\ell(y, z)$ some measure of how close a true label y is to a predicted label z .

⊕ **Example 2.1:**

- (Binary classification) Let $\mathcal{Y} = \{0, 1\}$, or even $\mathcal{Y} = \{0, \dots, 9\}$. A typical loss on such labels is the “0-1 loss”, $\ell(y, z) := \mathbb{1}_{\{y \neq z\}}$.
- (Regression) Let $\mathcal{Y} = \mathbb{R}$, then two typical loss functions are the *mean-square loss*

$$\ell(y, z) := (y - z)^2$$

or *absolute loss*

$$\ell(y, z) := |y - z|.$$

2.2.2 Risks

→**Definition 2.3** (Expected Risk): Given a prediction function $f : \mathcal{X} \rightarrow \mathcal{Y}$, a loss function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ and a probability distribution P on $\mathcal{X} \times \mathcal{Y}$, the *expected risk* of f is defined by

$$\mathcal{R}(f) := \mathbb{E}_{x,y}[\ell(y, f(x))] = \int_{\mathcal{X} \times \mathcal{Y}} \ell(y, f(x)) dP(x, y).$$

→**Definition 2.4** (Empirical Risk): Given a prediction function $f : \mathcal{X} \rightarrow \mathcal{Y}$, a loss function $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ and $(x_i, y_i)_{i=1}^n \in \mathcal{X} \times \mathcal{Y}$, the *empirical risk* is given by

$$\begin{aligned} \widehat{\mathcal{R}}(f) &:= \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)) \\ &= \int_{\mathcal{X} \times \mathcal{Y}} \ell(y, f(x)) d\mu(x, y), \quad \mu(x, y) := \frac{1}{n} \sum_{i=1}^n \delta_{\{(x_i, y_i)\}}. \end{aligned}$$

Remark 2.1: Heuristically, $\widehat{\mathcal{R}}(f)$ should approach $\mathcal{R}(f)$ as $n \rightarrow \infty$.

④ **Example 2.2:**

1. If $\mathcal{Y} = \{0, 1\}$, $\ell(y, z) = \mathbb{1}_{\{y \neq z\}}$, then

$$\mathcal{R}(f) = \mathbb{E}[\mathbb{1}_{\{y \neq f(x)\}}] = \mathbb{P}(y \neq f(x)) = \text{probability of missclassifying}$$

2. $\mathcal{Y} = \mathbb{R}$, $\ell(y, z) = (y - z)^2$,

$$\mathcal{R}(f) = \mathbb{E}[(y - f(x))^2], \quad \text{mean-square error (MSE)}$$

2.2.3 Baye's Risk, Predictor

Here, we answer the question: what's the best predictor f we could find, assuming we knew everything about the underlying distribution on $\mathcal{X} \times \mathcal{Y}$?

We can write, by law of total expectation,

$$\begin{aligned}\mathcal{R}(f) &= \mathbb{E}[\ell(y, f(x))] \\ &= \mathbb{E}[\mathbb{E}[\ell(y, f(x))|x]] \\ &= \mathbb{E}_{x' \sim p}[\mathbb{E}[\ell(y, f(x')) | x = x']] \\ &= \int_{\mathcal{X}} \mathbb{E}[\ell(y, f(x')) | x = x'] dp(x').\end{aligned}$$

Define the *conditional risk* given $x' \in \mathcal{X}$ by

$$r(z|x') := \mathbb{E}_y[\ell(y, z) | x = x'],$$

so that we can write

$$\mathcal{R}(f) = \int_{\mathcal{X}} r(f(x')|x) dp(x') \stackrel{\mathcal{X} \text{ finite}}{=} \sum_{x' \in \mathcal{X}} r(f(x')|x') \mathbb{P}(x = x').$$

In particular, in the finite case, we can see that to minimize the risk $\mathcal{R}(f)$, we can minimize the individual conditional risks $r(f(x')|x')$ for each $x' \in \mathcal{X}$. The so-called *Baye's predictor* is a function f_* which for each x' minimizes $r(f(x')|x')$. Formally,

→ **Proposition 2.1** (Baye's Predictor/Risk): The expected risk is minimized at a *Baye's predictor* $f_* : \mathcal{X} \rightarrow \mathcal{Y}$ such that, for all $x' \in \mathcal{X}$,

$$f_*(x') \in \operatorname{argmin}_{z \in \mathcal{Y}} \mathbb{E}[\ell(y, z) | x = x'].$$

The *Baye's risk* is the risk of a (any) Baye's predictor, written

$$\mathcal{R}^* := \mathbb{E}_{x' \sim p} \left[\inf_{z \in \mathcal{Y}} \mathbb{E}[\ell(y, z) | x = x'] \right] = \mathbb{E}[\ell(y, f_*(x'))].$$

Remark 2.2:

1. Finding an f_* is an impossible task in practice. Instead, we'll usually assume f takes some parametrized form, and optimize these parameters.
2. Baye's predictor may not be unique, but all Baye's predictors have the same risk
3. Baye's risk is usually nonzero, unless the dependency between x and y is deterministic.

→ **Definition 2.5** (Excess Risk): The *excess risk* of a predictor f is the value $\mathcal{R}(f) - \mathcal{R}(f_*) \geq 0$.

Remark 2.3: Thus, if we knew the conditional distribution $(y|x)$ for each x , the optimal predictor is known. ML can be succinctly be described as dealing with the general case in which we do not know $(y|x)$ for all x , and can only work with given samples of data.

⊕ **Example 2.3:**

1. (Binary classification) With $\mathcal{Y} := \{-1, 1\}$ and $\ell(y, z) = \mathbb{1}_{\{y \neq z\}}$ the 0-1 loss, we can see that

$$\begin{aligned} f_*(x') &\in \operatorname{argmin}_{z' \in \{-1, 1\}} P(y \neq z | x = x') \\ &= \operatorname{argmax}_{z \in \{-1, 1\}} \mathbb{P}(y = z | x = x') \\ &= \begin{cases} 1 & \mathbb{P}(y = 1 | x = x') > \frac{1}{2} \\ -1 & \mathbb{P}(y = 1 | x = x') < \frac{1}{2} \text{ anything } \mathbb{P}(y = 1 | x = x') = \frac{1}{2}. \end{cases} \end{aligned}$$

Putting $\mathcal{L}(x') := \mathbb{P}(y = 1 | x = x')$, this implies

$$\mathcal{R}^* = \mathbb{E}[\min\{\mathcal{L}(x), 1 - \mathcal{L}(x)\}].$$

2. (Regression) With $\mathcal{Y} = \mathbb{R}$, $\ell(y, z) = (y - z)^2$, we see that

$$\begin{aligned} \operatorname{argmin}_{z \in \mathbb{R}} \mathbb{E}[(y - z)^2 | x = x'] &= \operatorname{argmin}_{z \in \mathbb{R}} \left\{ \underbrace{\mathbb{E}[(y - \mathbb{E}[y | x = x'])^2 | x = x']}_{\text{independent of } z} \right. \\ &\quad \left. + \underbrace{(z - \mathbb{E}[y | x = x'])^2}_{\text{minimize this}} \right\} \\ &= \mathbb{E}[y | x = x']. \end{aligned}$$

Hence, $f_*(x') = \mathbb{E}[y | x = x']$, and so

$$\mathcal{R}^* = \mathbb{E}_{x' \sim p} \left[\inf_{z \in \mathbb{R}} \mathbb{E}[(y - z)^2 | x = x'] \right] = \mathbb{E}_{x'} [(y - \mathbb{E}[y | x = x'])^2] \quad (\text{conditional variance})$$

§2.3 Empirical Risk Minimization

We don't know the underlying distributions we work with (of course, otherwise we'd be done), and we need to work with samples, and need to simplify what kind of prediction functions we consider (since we don't know the underlying distribution, thus can't find the Baye's predictor in general).

We'll assume a parametrized family of predictor functions (called our *model*),

$$f_\theta : \mathcal{X} \rightarrow \mathcal{Y}, \quad \theta \in \Theta,$$

where $\Theta \subset \mathbb{R}^d$ typically. Heuristically, as d increases, if we could find the best f_θ predictor for $\theta \in \Theta$, that predictor will approach the Baye's predictor.

→ **Definition 2.6** (Empirical risk with respect to a parameter): The *empirical risk* w.r.t $\theta \in \Theta$ is

$$\hat{R}(f_\theta) := \frac{1}{n} \sum_{i=1}^n \ell(y_i, f_\theta(x_i)).$$

We consider the optimal parameter minimizing this empirical risk as

$$\hat{\theta} \in \operatorname{argmin}_{\theta \in \Theta} \hat{R}(f_\theta),$$

and so our “optimal” prediction function with respect to Θ is $f_{\hat{\theta}}$.

⊕ **Example 2.4:** A typical linear least-squares problem takes this form,

$$\min_{\theta \in \Theta = \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n (y_i - \theta^T \varphi(x_i))^2,$$

so that here, $f_\theta(x) = \theta^T \varphi(x)$ and our loss function is the square loss.

2.3.1 Risk Decomposition

Given a $\hat{\theta} \in \Theta$ (not necessarily optimal w.r.t Θ), we would like to break down the excess risk of the predictor $f_{\hat{\theta}}$ w.r.t the Baye’s predictor to see the difference in error coming from our choice of model (we call this *approximation error*, i.e. how far our model is from approximating our true predictor function) versus from the choice of $f_{\hat{\theta}}$ over the “true” best predictor with respect to Θ (as defined in the previous section). This is called the *estimation error*, and should be thought of as how well any underlying optimization algorithm used to find $\hat{\theta}$ performed compared to the theoretical best).

Mathematically, we can write

$$\underbrace{\mathcal{R}(f_{\hat{\theta}}) - \mathcal{R}^*}_{\text{Excess Risk}} = \underbrace{\left\{ \mathcal{R}(f_{\hat{\theta}}) - \inf_{\theta \in \Theta} \mathcal{R}(f_\theta) \right\}}_{\text{Estimation Error}} + \underbrace{\left\{ \inf_{\theta \in \Theta} \mathcal{R}(f_\theta) - \mathcal{R}^* \right\}}_{\text{Approximation Error}}.$$

how good our estimator is from the best possible how good our estimator is compared to the best the model can do how good our model (theoretically) is compared to the best possible

Note that the approximation error is due to the modelling choice, and is independent of the specific $f_{\hat{\theta}}$. Vaguely, “as Θ grows, the approximation error should shrink”.

The estimation error can further be broken down into three parts; let $\theta' \in \Theta$ be the minimizer of $\theta \mapsto \mathcal{R}(f_\theta)$ (e.g., $\mathcal{R}(f_{\theta'}) = \inf_{\theta \in \Theta} \mathcal{R}(f_\theta)$), then

$$\begin{aligned}
& \underbrace{\mathcal{R}(f_{\hat{\theta}}) - \mathcal{R}(f_{\theta'})}_{\text{Estimation Error}} = \{\mathcal{R}(f_{\hat{\theta}}) - \widehat{\mathcal{R}}(f_{\hat{\theta}})\} \leftarrow \text{How good the model risk is on data vs true risk of model} \\
& \text{Empirical Optimization Error} \\
& \text{How bad our choice of predictor is compared to the best in terms of performance on the data (for } \hat{\theta} \text{)} \rightarrow +\{\widehat{\mathcal{R}}(f_{\hat{\theta}}) - \widehat{\mathcal{R}}(f_{\theta'})\} \\
& +\{\widehat{\mathcal{R}}(f_{\theta'}) - \mathcal{R}(f_{\theta'})\} \leftarrow \text{How good the model risk is on data vs true risk of model (for } \theta' \text{)} \\
& \leq 2 \underbrace{\sup_{\theta \in \Theta} |\mathcal{R}(f_{\theta}) - \widehat{\mathcal{R}}(f_{\theta})|}_{\text{should } \downarrow \text{ as } n \uparrow} + \underbrace{\{\widehat{\mathcal{R}}(f_{\hat{\theta}}) - \widehat{\mathcal{R}}(f_{\theta'})\}}_{\substack{\uparrow \text{ as } \Theta \uparrow \\ (\Theta \text{ gets too large to optimize over})}} .
\end{aligned}$$

In brief, we expect that as the parameter space Θ grows, the estimation error *increases*, but the approximation error *decreases*. But as n (number of samples) increases, we expect the estimation error to decrease (and there is no effect on the approximation error). Thus, there is a subtle interplay between $d := \dim(\Theta)$ and n .

§2.4 Statistical Learning Theory

“Statistical learning theory” asks how to provide guarantees of performance of an algorithm on previously unseen data.

We assume we have data

$$D_n(p) := \{(x_1, y_1), \dots, (x_n, y_n)\}$$

which are assumed to be iid from some unknown distribution p which is part of some family P of distributions.

An algorithm then is a mapping A from $D_n(p)$ to a prediction function $A(D_n(p)) = f : \mathcal{X} \rightarrow \mathcal{Y}$. Our goal is to find an algorithm such that the excess risk of the prediction function given by A is “small”, in a sense we’ll define in the next section.

2.4.1 Measures of Performance

→ **Definition 2.7** (Expected Risk): The *expected risk* of an algorithm A on sample size n and probability distribution p is the quantity

$$\mathbb{E}[\mathcal{R}_p(A(D_n(p)))],$$

where the expected value is taken over all possible n -size subsets of the training data. We say that an algorithm is *consistent in expectation* if the above quantity converges, with p fixed, to \mathcal{R}^* as $n \rightarrow \infty$.

→ **Definition 2.8** (Probability Approximately Correct *): We say an algorithm A is Probability Approximately Correct (PAC) if for any given $\delta \in (0, 1)$ and $\varepsilon > 0$, there exists $N \in \mathbb{N}$ such that for all $n \geq N$,

$$\mathbb{P}(\mathcal{R}_p(A(D_n(p))) - \mathcal{R}^* \leq \varepsilon) \geq 1 - \delta.$$

2.4.2 Notions of Consistency over Classes of Probability distributions

Remark that our definition of consistency in expectation gave no guarantee over rate of convergence, especially not with respect to the specific distribution.

→**Definition 2.9:** An algorithm is *uniformly consistent* if for all probability distributions on (x, y) , the algorithm is consistent.

→**Definition 2.10** (Minimax risk): The minimax risk is defined to be, given $\mathcal{X} \times \mathcal{Y}$,

$$\inf_{A: \text{ algorithm}} \sup_{\substack{p \in P: \\ \text{class of dists.}}} \left\{ \mathbb{E}[\mathcal{R}_p(A(D_n(p)))] - \mathcal{R}^* \right\}.$$

Remark 2.4: This is hard to evaluate in general, but is easy to upper bound (just fix any A and evaluate the inner supremum, i.e., look at the worst-case performance of the algorithm). Lower bounds are much harder to compute, since they need to hold for any possible algorithm.